



Troubleshooting & Debugging


Chirag Shankar

Systems Engineer, Google Cloud




Troubleshooting & Debug

Agenda




Troubleshooting Workflow



Troubleshooting
workflow

- Checking for errors
- Looking for anomalies in the Job metrics tab


Troubleshooting Workflow



Troubleshooting workflow

- Checking for errors
 - Logs
 - Execution Graph
- Looking for anomalies in the Job metrics tab


Checking for errors




Dataflow





Jobs

[+ CREATE JOB FROM TEMPLATE](#)



Filter jobs



Name	Type	End time	Elapsed time	Start time	Status	SDK version	ID	Region
 wordcount4	Batch	—	7 sec	Nov 21, 2018, 3:22:45 PM	Running	2.7.0	2018-11-21_15_22_45-	us-central1
 wordcount3	Batch	Nov 21, 2018, 1:10:23 PM	3 min 29 sec	Nov 21, 2018, 1:06:54 PM	Succeeded	2.7.0	2018-11-21_13_06_53-	us-central1
 wordcount2	Batch	Nov 21, 2018, 12:57:05 PM	3 min 56 sec	Nov 21, 2018, 12:53:09 PM	Failed	2.7.0	2018-11-21_12_53_08-	us-central1
 wordcount1	Batch	Nov 21, 2018, 12:55:25 PM	3 min 36 sec	Nov 21, 2018, 12:51:49 PM	Succeeded	2.7.0	2018-11-21_12_51_48-	us-central1

Checking for errors

Dataflow

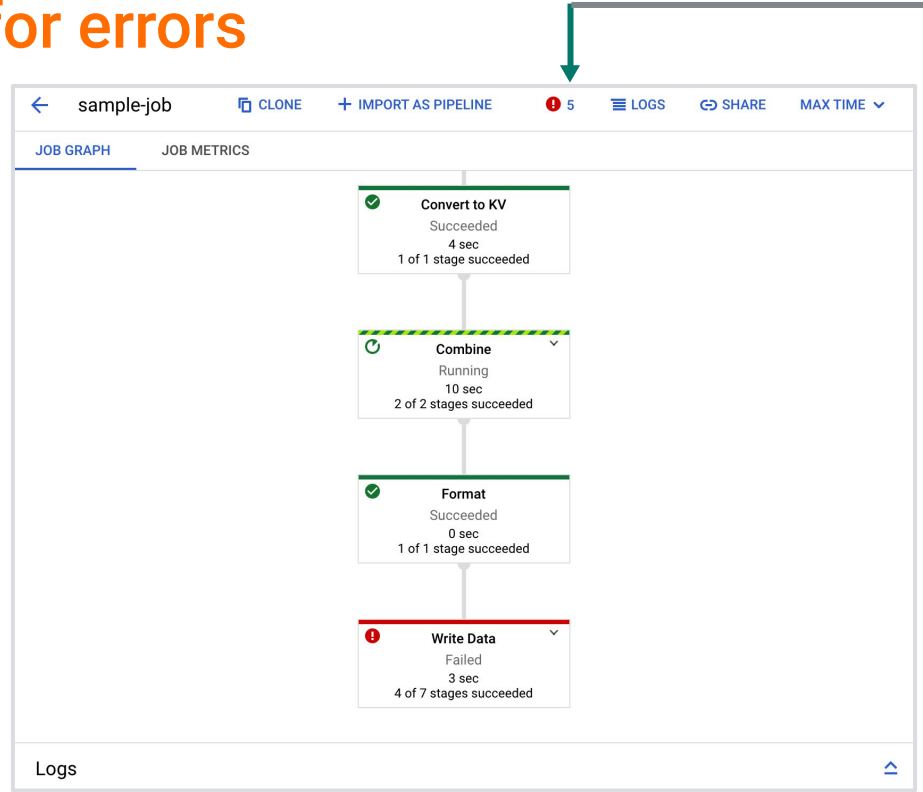
Jobs

+ CREATE JOB FROM TEMPLATE

Filter jobs

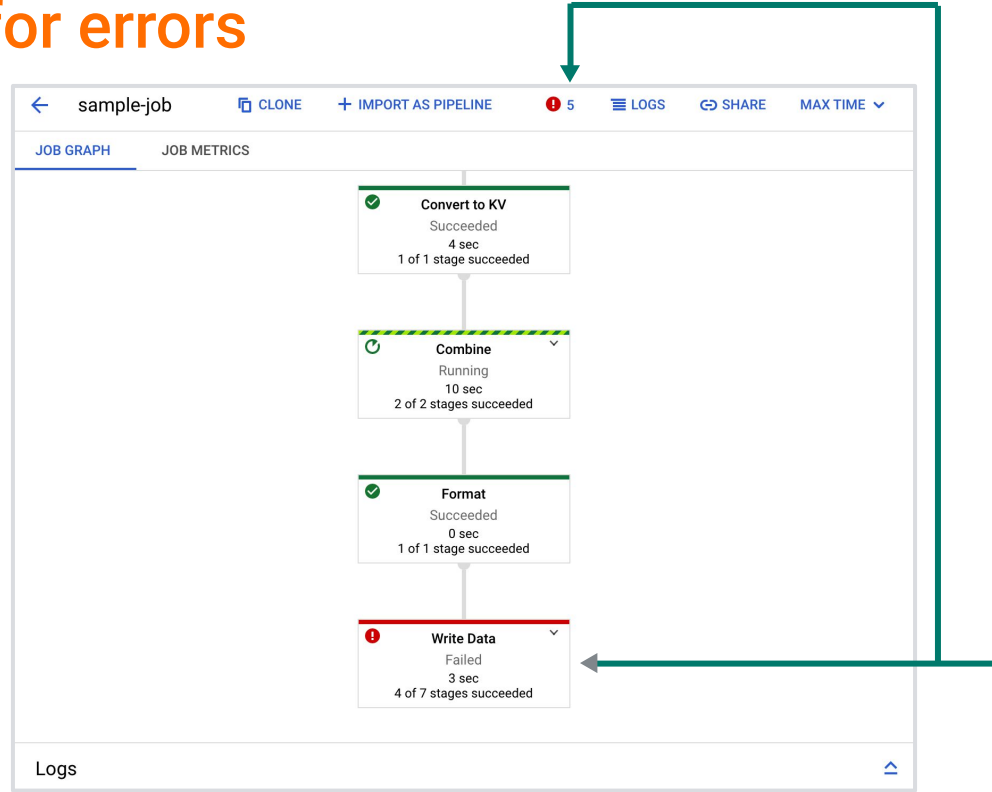
Name	Type	End time	Elapsed time	Start time	Status	SDK version	ID	Region
<div><div></div><div>wordcount4</div></div>	Batch	—	7 sec	Nov 21, 2018, 3:22:45 PM	Running	2.7.0	2018-11-21_15_22_45-2018-11-21-15-22-45-2018-11-21-15-22-45	us-central1
<div><div></div><div>wordcount3</div></div>	Batch	Nov 21, 2018, 1:10:23 PM	3 min 29 sec	Nov 21, 2018, 1:06:54 PM	Succeeded	2.7.0	2018-11-21_13_06_53-2018-11-21-13-06-53-2018-11-21-13-06-53	us-central1
<div><div></div><div>wordcount2</div></div>	Batch	Nov 21, 2018, 12:57:05 PM	3 min 56 sec	Nov 21, 2018, 12:53:09 PM	Failed	2.7.0	2018-11-21_12_53_08-2018-11-21-12-53-08-2018-11-21-12-53-08	us-central1
<div><div></div><div>wordcount1</div></div>	Batch	Nov 21, 2018, 12:55:25 PM	3 min 36 sec	Nov 21, 2018, 12:51:49 PM	Succeeded	2.7.0	2018-11-21_12_51_48-2018-11-21-12-51-48-2018-11-21-12-51-48	us-central1

Checking for errors



Checking for errors

Checking for errors



Checking
for errors

Checking for errors

The screenshot displays the Google Cloud Data Studio interface for a job named "sample-job". At the top, there are navigation links for "CLONE", "LOGS", "SHARE", and "MAX TIME". Below the job name, there are tabs for "JOB GRAPH" and "JOB METRICS". The "JOB GRAPH" tab is active, showing a single stage "1 of 1 stage succeeded". A red box highlights a "Write Data" operation that failed, with the message "Failed 3 sec 4 of 7 stages succeeded". Below the job graph, there is a "Logs" section with tabs for "JOB LOGS", "WORKER LOGS", and "DIAGNOSTICS". The "JOB LOGS" tab is active, showing a list of 50 messages. The messages include timestamps, log levels, and descriptions of the job's progress and errors.

sample-job CLONE 5 LOGS SHARE MAX TIME

JOB GRAPH JOB METRICS

1 of 1 stage succeeded

Write Data
Failed
3 sec
4 of 7 stages succeeded

Logs

JOB LOGS WORKER LOGS DIAGNOSTICS

Showing 50 messages Info Filter Filter logs

- 2021-02-16 13:48:02.861 CST Executing operation Write Data/BigQueryBatchFileLoads/WaitForDestinationLoadJobs/WaitForDes...
- 2021-02-16 13:48:05.261 CST Error message from worker: Traceback (most recent call last): File "/usr/local/lib/python2...
- 2021-02-16 13:48:05.906 CST Error message from worker: Traceback (most recent call last): File "/usr/local/lib/python2...
- 2021-02-16 13:48:07.603 CST Error message from worker: Traceback (most recent call last): File "/usr/local/lib/python2...
- 2021-02-16 13:48:10.354 CST Error message from worker: Traceback (most recent call last): File "/usr/local/lib/python2...
- 2021-02-16 13:48:10.714 CST Finished operation Write Data/BigQueryBatchFileLoads/WaitForDestinationLoadJobs/WaitForDest...
- 2021-02-16 13:48:10.797 CST Workflow failed. Causes: S19:Write Data/BigQueryBatchFileLoads/WaitForDestinationLoadJobs/W...
- 2021-02-16 13:48:11.061 CST Stopping worker pool...
- 2021-02-16 13:50:15.803 CST Worker pool stopped.
- No newer entries found matching current filter.

Checking for errors

The screenshot displays the Google Cloud Dataflow console for a job named 'sample-job'. The job status is 'Failed', indicated by a red exclamation mark icon and the text '1 of 1 stage succeeded'. A detailed view of the failed stage, 'Write Data', shows it 'Failed' after '3 sec' and that '4 of 7 stages succeeded'. Below the job graph, the 'Logs' section is expanded, showing 'JOB LOGS', 'WORKER LOGS', and 'DIAGNOSTICS'. The 'JOB LOGS' tab is active, displaying a list of 50 messages. A red arrow points to the 'Filter' button in the log header, which is used to filter the log entries. The log entries show a sequence of events: executing the 'Write Data' operation, receiving error messages from workers (Traceback), finishing the operation, and finally failing the workflow due to an S19 error. The log also shows the worker pool being stopped.

sample-job CLONE 5 LOGS SHARE MAX TIME

JOB GRAPH JOB METRICS

1 of 1 stage succeeded

Write Data
Failed
3 sec
4 of 7 stages succeeded

Logs

JOB LOGS WORKER LOGS DIAGNOSTICS

Showing 50 messages Info Filter Filter logs

- 2021-02-16 13:48:02.861 CST Executing operation Write Data/BigQueryBatchFileLoads/WaitForDestinationLoadJobs/WaitForDes...
- 2021-02-16 13:48:05.261 CST Error message from worker: Traceback (most recent call last): File "/usr/local/lib/python2...
- 2021-02-16 13:48:05.906 CST Error message from worker: Traceback (most recent call last): File "/usr/local/lib/python2...
- 2021-02-16 13:48:07.603 CST Error message from worker: Traceback (most recent call last): File "/usr/local/lib/python2...
- 2021-02-16 13:48:10.354 CST Error message from worker: Traceback (most recent call last): File "/usr/local/lib/python2...
- 2021-02-16 13:48:10.714 CST Finished operation Write Data/BigQueryBatchFileLoads/WaitForDestinationLoadJobs/WaitForDest...
- 2021-02-16 13:48:10.797 CST Workflow failed. Causes: S19:Write Data/BigQueryBatchFileLoads/WaitForDestinationLoadJobs/W...
- 2021-02-16 13:48:11.061 CST Stopping worker pool...
- 2021-02-16 13:50:15.803 CST Worker pool stopped.
- No newer entries found matching current filter.

Checking for errors


The screenshot shows the Google Cloud Logs Explorer interface. At the top, there's a navigation bar with 'Logs Explorer', 'OPTIONS', 'REFINE SCOPE' (set to 'Project'), 'SHARE LINK', a time range of '1:44:48 PM - 1:51:27 PM', 'PAGE LAYOUT', and 'LEARN'. Below this is a 'Query preview' section with the query: `resource.type="dataflow_step" resource.labels.job_id="2021-02-16_11_44_47-5044811639078349390" logNam...`. There are buttons for 'Save', 'Stream logs', and 'Run query'.

On the left, a 'Log fields' sidebar is open, showing a search bar and a list of fields. The 'SEVERITY' section is expanded, showing 'Debug' (83), 'Info' (52), and 'Error' (5). The 'LOG NAME' section is also expanded, showing 'dataflow.googleapis.com/job-message' (140). A red arrow points from this field to the 'Query results' table.

The 'Query results' table has columns for 'SEVERITY', 'TIMESTAMP', 'CST', and 'SUMMARY'. It lists several error messages from workers, including 'Error message from worker: Traceback (most recent call last): File "/usr/...' and 'Workflow failed. Causes: S19:Write Data/BigQueryBatchFileLoads/WaitForDes...'. A red arrow points from the 'dataflow.googleapis.com/job-message' field in the sidebar to the first row of the table.

On the right, a 'Histogram' chart shows the distribution of log entries over time. A 'View details' button is visible. Below the histogram, there's a 'Work' section with a list of items: 'Job-', 'Work', 'Work', 'Shuf', and 'Dock'.

Troubleshooting Workflow




Troubleshooting workflow

- Checking for errors
- Looking for anomalies in the Job metrics tab

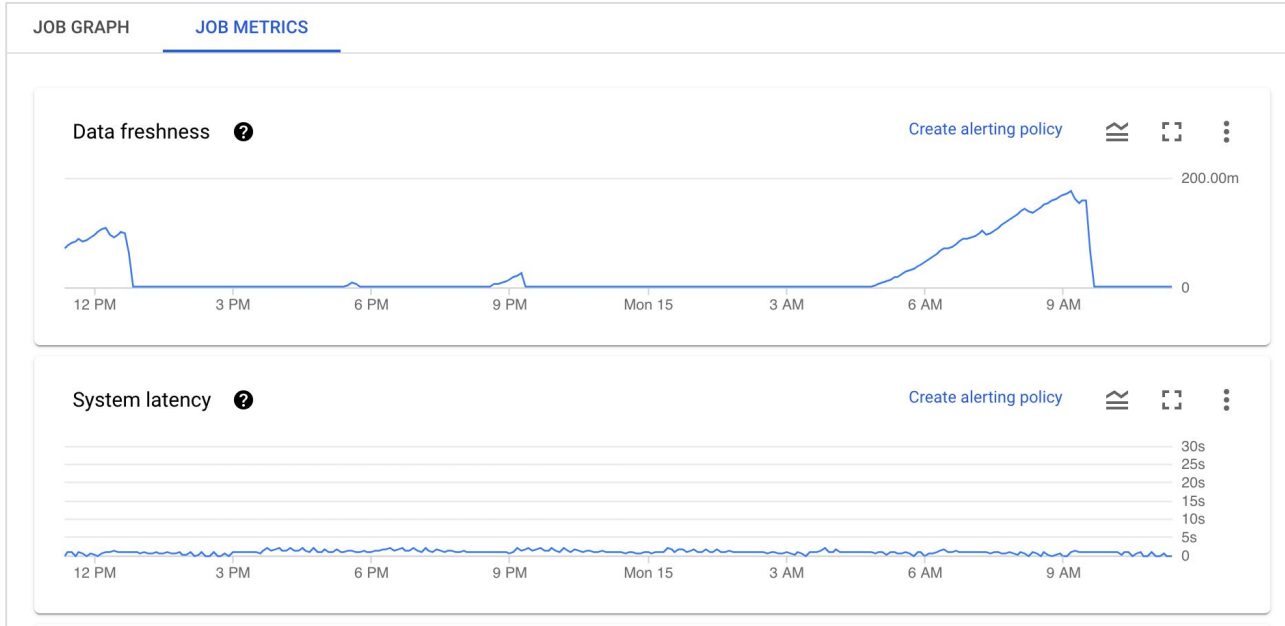
Looking for anomalies in the Job metrics tab

Streaming pipelines



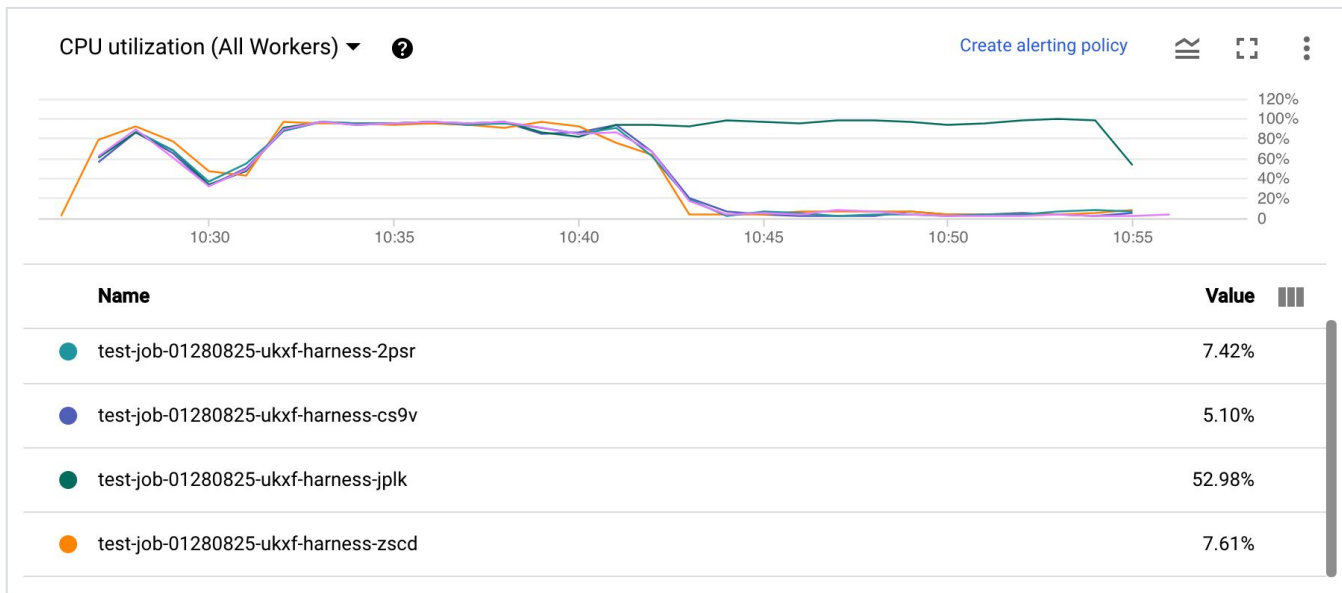
Looking for anomalies in the Job metrics tab

Streaming pipelines




Looking for anomalies in the Job metrics tab

All Dataflow Jobs




Troubleshooting & Debug

Agenda




Types of Troubles



Types of
troubles

- Failure building the pipeline
- Failure starting the pipeline on Dataflow
- Failure during pipeline execution
- Performance problems

Types of Troubles



Types of
troubles

- Failure building the pipeline
- Failure starting the pipeline on Dataflow
- Failure during pipeline execution
- Performance problems

Failure building the pipeline


These errors occur when Apache Beam is building the pipeline:

- Validating "Beam Model" aspects of the pipeline
- Validating input/output specifications
- You can reproduce it with the direct runner, in a unit test

Failure building the pipeline

```
java.lang.IllegalStateException: GroupByKey cannot be applied to non-bounded
PCollection in the GlobalWindow without a trigger. Use a Window.into or
Window.triggering transform prior to GroupByKey.
    at org.apache.beam.sdk.transforms.GroupByKey.applicableTo (GroupByKey.java:153)
    at org.apache.beam.sdk.transforms.GroupByKey.expand (GroupByKey.java:185)
    at org.apache.beam.sdk.transforms.GroupByKey.expand (GroupByKey.java:107)
    at org.apache.beam.sdk.Pipeline.applyInternal (Pipeline.java:537)
    at org.apache.beam.sdk.Pipeline.applyTransform (Pipeline.java:471)
    at org.apache.beam.sdk.values.PCollection.apply (PCollection.java:357)
    at ...
```

Types of Troubles



Types of
troubles

- Failure building the pipeline
- Failure starting the pipeline on Dataflow
- Failure during pipeline execution
- Performance problems

Failure starting the pipeline on Dataflow

Once the Dataflow service has received your pipeline's graph, the service will attempt to validate your job:

- Making sure the service can access your job's associated Cloud Storage buckets for file staging and temporary output.



Failure starting the pipeline on Dataflow

Once the Dataflow service has received your pipeline's graph, the service will attempt to validate your job:

- Making sure the service can access your job's associated Cloud Storage buckets for file staging and temporary output.
- Checking for the required permissions in your Google Cloud project.



Failure starting the pipeline on Dataflow

Once the Dataflow service has received your pipeline's graph, the service will attempt to validate your job:


- Making sure the service can access your job's associated Cloud Storage buckets for file staging and temporary output.
- Checking for the required permissions in your Google Cloud project.
- Making sure the service can access input and output sources, such as files.



Failure starting the pipeline on Dataflow

```
apitools.base.py.exceptions.HttpForbiddenError: HttpError accessing
<https://dataflow.googleapis.com/v1b3/projects/xyz/locations/us-central1/jobs?alt=json>:
response: <{'status': '403', 'content-length': '288', 'x-xss-protection': '0',
'x-content-type-options': 'nosniff', 'transfer-encoding': 'chunked', 'vary': 'Origin,
X-Origin, Referer', 'server': 'ESF', '-content-encoding': 'gzip', 'cache-control': 'private',
'date': 'Tue, 10 Dec 2019 17:56:00 GMT', 'x-frame-options': 'SAMEORIGIN', 'alt-svc':
'quic=":443"; ma=2592000; v="46,43",h3-Q050=":443"; ma=2592000,h3-Q049=":443";
ma=2592000,h3-Q048=":443"; ma=2592000,h3-Q046=":443"; ma=2592000,h3-Q043=":443"; ma=2592000',
'content-type': 'application/json; charset=UTF-8'}>, content <{
  "error": {
    "code": 403,
    "message": "(43abf11c2446750): Could not create workflow; user does not have write
access to project: xyz Causes: (43abf11c244683b): Permission 'dataflow.jobs.create'
denied on project: 'xyz'",
    "status": "PERMISSION_DENIED"
  }
}
```

Types of Troubles



Types of
troubles

- Failure building the pipeline
- Failure starting the pipeline on Dataflow
- Failure during pipeline execution
- Performance problems

Failure during pipeline execution

Logs

JOB LOGS **WORKER LOGS** ⚠️ DIAGNOSTICS

Showing 78 messages

Severity
Error

Filter Filter logs

🔄 📄 ⓘ

!!	2021-03-08 11:30:53.696 CST	Uncaught exception:
!!	2021-03-08 11:30:53.696 CST	Uncaught exception:
!!	2021-03-08 11:30:53.697 CST	Uncaught exception:
!!	2021-03-08 11:30:53.698 CST	Uncaught exception:
!!	2021-03-08 11:30:53.698 CST	Uncaught exception:
!!	2021-03-08 11:30:53.701 CST	Uncaught exception:
!!	2021-03-08 11:30:53.704 CST	Uncaught exception:
!!	2021-03-08 11:30:53.707 CST	Uncaught exception:

Failure during pipeline execution

Log fields

Search fields and values

RESOURCE TYPE

Dataflow Step

SEVERITY

Error

LOG NAME

dataflow.googleapis.com/worker

PROJECT ID

JOB ID

STEP ID

(Missing or invalid value)

JOB NAME

REGION

us-central1

Query results

SEVERITY

2021-03-08 11:22:40.458 CST

Uncaught exception: "

2021-03-08 11:22:40.458 CST

Uncaught exception: "

Hide log summary

Collapse nested fields

Copy to clipboard

Copy link

```
{
  insertId: "776189048441790998:16632:0:753343"
  jsonPayload: {
    exception:
      "java.lang.RuntimeException: org.apache.beam.sdk.util.UserCodeException: java.lang.OutOfMemoryError: unable to create native thread: possibly out of memory or process/resource limits reached
        at org.apache.beam.runners.dataflow.worker.GroupAlsoByWindowsParDoFn$1.output(GroupAlsoByWindowsParDoFn.java:184)
        at org.apache.beam.runners.dataflow.worker.GroupAlsoByWindowFnRunner$1.outputWindowedValue(GroupAlsoByWindowFnRunner.java:107)
        at org.apache.beam.runners.dataflow.worker.repackaged.org.apache.beam.runners.core.ReduceFnRunner.lambda$onTrigger$1(ReduceFnRunner.java:1057)
        at org.apache.beam.runners.dataflow.worker.repackaged.org.apache.beam.runners.core.ReduceFnContextFactory$onTriggerContextImpl.output(ReduceFnContextFactory.java:438)
        at org.apache.beam.runners.dataflow.worker.repackaged.org.apache.beam.runners.core.SystemReduceFn.onTrigger(SystemReduceFn.java:125)
        at org.apache.beam.runners.dataflow.worker.repackaged.org.apache.beam.runners.core.ReduceFnRunner.onTrigger(ReduceFnRunner.java:1060)
        at org.apache.beam.runners.dataflow.worker.repackaged.org.apache.beam.runners.core.ReduceFnRunner.onTimers(ReduceFnRunner.java:768)
        at org.apache.beam.runners.dataflow.worker.StreamingGroupAlsoByWindowViaWindowSetFn.processElement(StreamingGroupAlsoByWindowViaWindowSetFn.java:95)
        at org.apache.beam.runners.dataflow.worker.StreamingGroupAlsoByWindowViaWindowSetFn.processElement(StreamingGroupAlsoByWindowViaWindowSetFn.java:42)
        at org.apache.beam.runners.dataflow.worker.GroupAlsoByWindowFnRunner.invokeProcessElement(GroupAlsoByWindowFnRunner.java:120)
        at org.apache.beam.runners.dataflow.worker.GroupAlsoByWindowFnRunner.processElement(GroupAlsoByWindowFnRunner.java:73)
        at org.apache.beam.runners.dataflow.worker.repackaged.org.apache.beam.runners.core.LateDataDroppingDoFnRunner.processElement(LateDataDroppingDoFnRunner.java:80)
        at org.apache.beam.runners.dataflow.worker.GroupAlsoByWindowsParDoFn.processElement(GroupAlsoByWindowsParDoFn.java:134)
        at org.apache.beam.runners.dataflow.worker.util.common.worker.ParDoOperation.process(ParDoOperation.java:44)
        at org.apache.beam.runners.dataflow.worker.util.common.worker.OutputReceiver.process(OutputReceiver.java:49)
        at org.apache.beam.runners.dataflow.worker.util.common.worker.ReadOperation.runReadLoop(ReadOperation.java:201)
        at org.apache.beam.runners.dataflow.worker.util.common.worker.ReadOperation.start(ReadOperation.java:159)
        at org.apache.beam.runners.dataflow.worker.util.common.worker.MapTaskExecutor.execute(MapTaskExecutor.java:77)
        at org.apache.beam.runners.dataflow.worker.StreamingDataflowWorker.process(StreamingDataflowWorker.java:1363)
        at org.apache.beam.runners.dataflow.worker.StreamingDataflowWorker.access$1100(StreamingDataflowWorker.java:153)
        at org.apache.beam.runners.dataflow.worker.StreamingDataflowWorker$7.run(StreamingDataflowWorker.java:1086)
        at java.base/java.util.concurrent.ThreadPoolExecutor.runWorker(ThreadPoolExecutor.java:1128)
        at java.base/java.util.concurrent.ThreadPoolExecutor$Worker.run(ThreadPoolExecutor.java:628)
        at java.base/java.lang.Thread.run(Thread.java:834)
      "
  }
}
```

Caused by: org.apache.beam.sdk.util.UserCodeException: java.lang.OutOfMemoryError: unable to create native thread: possibly out of memory or process/resource limits reached

Beam College 2021

Failure during pipeline execution

```
class FilterMessagesFn(beam.DoFn):
    BAD_MESSAGE_TAG = 'bad_message'
    GOOD_MESSAGE_TAG = 'good_message'

    def process(self, element, window=beam.DoFn.WindowParam):
        try:
            data = element.decode()
            # tag the elements accordingly
            if 'bad' in data:
                yield pvalue.TaggedOutput(self.BAD_MESSAGE_TAG, element)
            else:
                yield pvalue.TaggedOutput(self.GOOD_MESSAGE_TAG, element)

        # handle any exceptions in the processing
        except Exception as exp:
            logging.getLogger.warning(exp)
            yield pvalue.TaggedOutput(self.BAD_MESSAGE_TAG, element)
```

Failure during pipeline execution

To track failing elements:

- Log the failing elements and check the output using Cloud Logging.



Failure during pipeline execution

To track failing elements:

- Log the failing elements and check the output using Cloud Logging.
- Check the Dataflow worker and worker startup logs for warnings or errors by following the instructions in Viewing logs.



Failure during pipeline execution


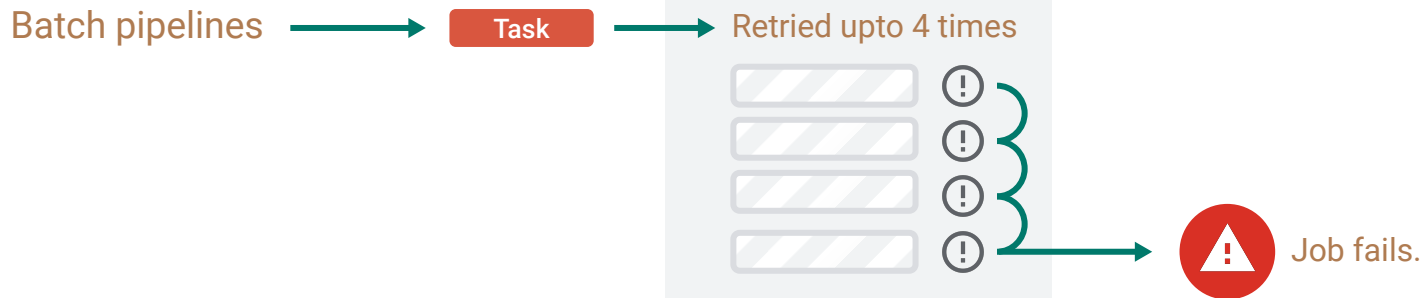
To track failing elements:

- Log the failing elements and check the output using Cloud Logging.
- Check the Dataflow worker and worker startup logs for warnings or errors by following the instructions in Viewing logs.
- Set your ParDo to write the failing elements to an additional output for later inspection.




Failure during pipeline execution

Batch vs Streaming



Types of Troubles



Types of
troubles

- Failure building the pipeline
- Failure starting the pipeline on Dataflow
- Failure during pipeline execution
- Performance problems

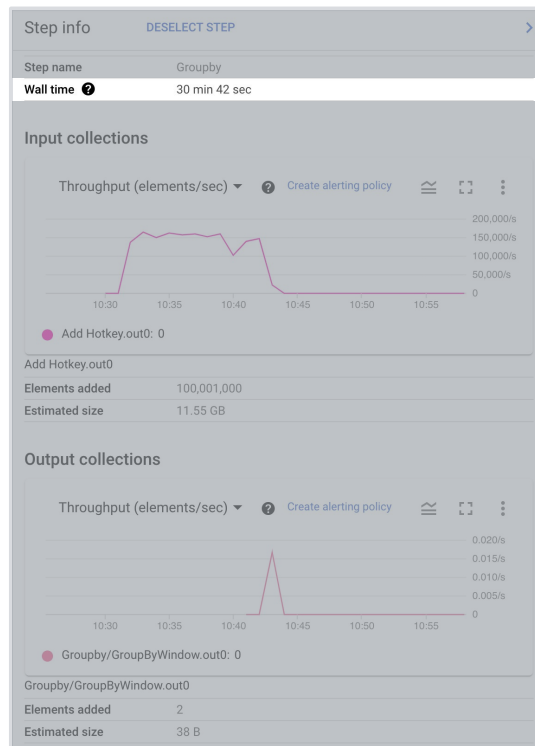
Performance problems

Factors that can influence the performance of a pipeline:

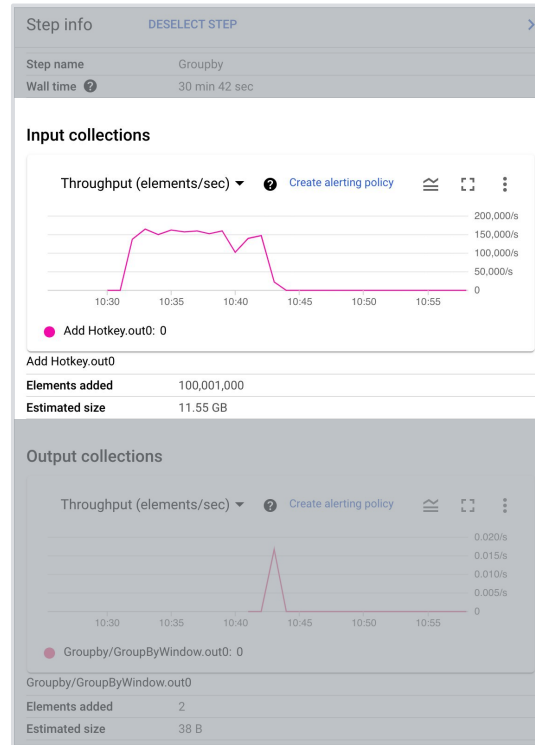
- Pipeline design
- Data shape
- Interactions with sources, sinks and external systems



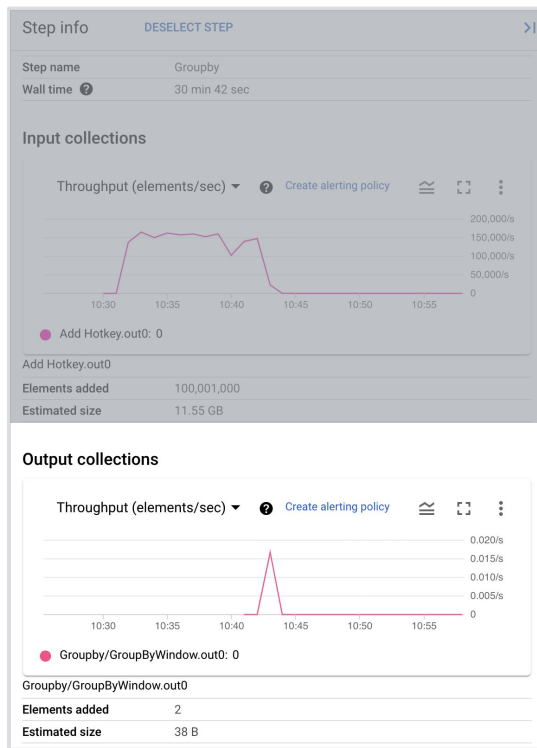
Performance problems



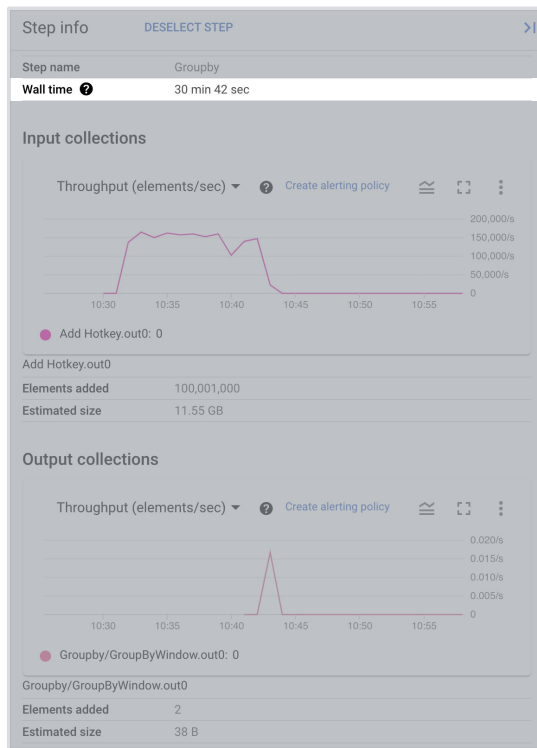
Performance problems




Performance problems



Performance problems



Performance problems



Thank you!

Questions?

